


Canvas Medical

Documentation Infrastructure Audit Report

Prepared by VeriOps

2026-04-25 13:24 UTC

CONFIDENTIAL | FOR INTERNAL USE ONLY



Risk: Moderate

Key Metrics

Pages scanned: 787

Report type: Premium Executive Audit


Methodology: 5-pillar automated analysis + expert synthesis

Headline Findings

- Broken internal links: 36
- API usage-doc coverage is low: 4.95% (273 API pages detected)

This report is generated automatically by the VeriOps platform. Estimates are directional.

Executive Summary



Documentation audit of docs.canvasmedical.com covered 787 pages with 100% crawl coverage across 790 discovered pages. The site contains 36 confirmed broken internal links affecting user-facing documentation. API reference coverage is critically low at 4.95% despite 273 API pages detected, indicating minimal integration of usage examples with API documentation. Example code reliability is estimated at 80%, and freshness metadata coverage is strong at 99.49%. Overall audit confidence stands at 91.5%.

External posture: SEO/GEO issue rate **7.9%**, public API coverage **5.0%**, broken links **36**.

73 Audit Score	B Grade	Moderate Risk Band
787 Pages Scanned	36 Broken Links	7.9% SEO Issue Rate

Per-Site Breakdown


Site URL	Pages	Broken Links	API Cov %	Example Rel %	SEO Issue %
https://docs.canvasmedical.com/	787	36	5.0	80.0	7.9

Industry Benchmark Comparison


Benchmark	Score	Gap
Top performers (Stripe, Twilio)	92+	-19
Industry average (developer docs)	85	-12
Minimum acceptable	70	+3
Your score	73	-

Gap to industry average: **12 points**. Top-performing developer documentation platforms (Stripe, Twilio, Vercel) score 92+. Closing this gap reduces support ticket volume by an estimated 20-30%.

Board-Level Metrics



Internal Metrics (from scorecard)



Financial Exposure Model

Low Estimate	Base Estimate	High Estimate
\$140,822 annual exposure	\$249,513 annual exposure	\$431,728 annual exposure
Remediation: \$2,090 Monthly loss: \$4,125 Opportunity: \$7,436/mo	Remediation: \$5,445 Monthly loss: \$12,903 Opportunity: \$7,436/mo	Remediation: \$8,800 Monthly loss: \$27,808 Opportunity: \$7,436/mo

Remediation: finding-level effort model | Monthly loss: operational friction model | Opportunity: engineering/support/release delay

Risk Matrix

ID	Issue	Severity	Monthly Loss	Fix Cost
F-LAYERS	Missing required doc layers	HIGH	\$1,782	\$825
F-FRESHNESS	Documentation freshness debt	MEDIUM	\$1,122	\$550
F-RETRIEVAL	RAG retrieval quality risk	MEDIUM	\$2,257	\$990
F-API-COVERAGE	Undocumented API operations	LOW	\$2,904	\$1,100
F-EXAMPLES-RELIABILITY	Code examples fail or are non-runnable	LOW	\$2,244	\$935
F-DRIFT	Code/docs contract drift	LOW	\$1,901	\$660
F-TERMINOLOGY	Terminology inconsistency	LOW	\$693	\$385

Priority Action Plan (Next 14 Days)

ID	Finding	Severity
F-LAYERS	Missing required doc layers	HIGH
F-FRESHNESS	Documentation freshness debt	MEDIUM
F-RETRIEVAL	RAG retrieval quality risk	MEDIUM
F-API-COVERAGE	Undocumented API operations	LOW
F-EXAMPLES-RELIABILITY	Code examples fail or are non-runnable	LOW

Recommended Actions

- Fix all 36 broken internal links in user-facing documentation [impact: critical, effort: low]
- Integrate usage examples and tutorials into 273 detected API pages to raise coverage from 4.95% to minimum 40% [impact: critical, effort: high]

3. Audit and update the 20% of code examples flagged as potentially unreliable [impact: high, effort: medium]
4. Resolve SEO/geo issues affecting 7.91% of pages to improve search discoverability [impact: medium, effort: medium]
5. Investigate 318 trap URLs to eliminate crawl inefficiencies and ensure complete indexing [impact: medium, effort: low]

Expert Analysis: Strengths and Risks

Strengths

- + Complete crawl coverage at 100% with 787 pages successfully indexed
- + High freshness metadata coverage at 99.49% enables content currency tracking
- + Strong overall audit confidence at 91.5% with 95% crawl confidence
- + Single-product topology provides clear documentation structure
- + No repository navigation link issues (0 repo broken links)

Risks

- 36 confirmed broken internal links in user-facing documentation degrade user experience and trust
- Critically low API reference coverage at 4.95% leaves developers without adequate usage guidance
- 273 API pages detected but minimal integration with examples creates documentation gaps
- SEO/geo issue rate of 7.91% may impact discoverability and search rankings
- Example reliability at 80% means 1 in 5 code samples may be outdated or non-functional
- 318 trap URLs skipped during crawl suggests potential pagination or infinite scroll issues
- API confidence at 85% indicates moderate uncertainty in API documentation assessment

Limitations

- * External audit cannot verify accuracy of medical/clinical content or regulatory compliance requirements
- * Code example functionality testing requires execution environment; audit provides reliability estimates only
- * API coverage metrics depend on detection heuristics; actual developer experience may vary
- * Broken link detection is point-in-time; dynamic content or authentication-gated pages may not be fully assessed
- * User comprehension, workflow completeness, and domain-specific terminology appropriateness require subject matter expert review

Methodology

This audit applies a rigorous, repeatable methodology combining automated analysis with expert synthesis. The platform evaluates documentation quality across five pillars.

1. Automated Crawl + Structural HTML Analysis

Every public documentation page is crawled and analyzed for structural integrity, link health, navigation depth, and content organization. The crawler detects broken links, orphaned pages, redirect chains, and missing metadata.

2. SEO/GEO Scoring Model (24 Checks)

Each page is evaluated against 8 GEO checks (LLM/AI search optimization) and 16 SEO checks (traditional search engine optimization). Checks cover meta descriptions, heading hierarchy, fact density, URL structure, internal linking, and content depth.

3. API Reference Cross-Coverage Analysis

The platform compares published API reference documentation against the actual API surface (OpenAPI specs, GraphQL schemas, gRPC protos). Coverage gaps, undocumented endpoints, and stale references are identified.

4. Code Example Reliability Estimation

Code examples embedded in documentation are analyzed for syntactic correctness, completeness, and consistency with current API signatures. Reliability scores estimate the likelihood of examples working as documented.

5. Executive Synthesis and Prioritization

All quantitative findings are synthesized into actionable executive narratives. The analysis identifies patterns, prioritizes risks, and generates strategic recommendations calibrated to business impact.

Appendix A -- Economic Model Assumptions

Assumption	Value
engineer_hourly_usd	110.0
support_hourly_usd	50.0
release_count_per_month	8.0
baseline_manual_sync_hours_per_week	12.0
avg_release_delay_hours	4.5
monthly_support_tickets	800.0
docs_related_ticket_share	0.22
avg_ticket_handling_minutes	24.0

Note: estimates are directional and should be calibrated with client rates, release cadence, and support volume.

Appendix B -- Evidence Traceability

Finding	Evidence Source	Confidence
F-LAYERS -- Missing required doc layers	policy pack + frontmatter layer scan	HIGH
F-FRESHNESS -- Documentation freshness debt	frontmatter date scan	HIGH
F-RETRIEVAL -- RAG retrieval quality risk	retrieval_evals_report.json	HIGH
F-API-COVERAGE -- Undocumented API operations	OpenAPI + docs text scan	HIGH
F-EXAMPLES-RELIABILITY -- Code examples fail or are non-runnable	examples_smoke_report.json	HIGH
F-DRIFT -- Code/docs contract drift	pr_docs_contract.json + api_sdk_drift_report.json	HIGH
F-TERMINOLOGY -- Terminology inconsistency	glossary.yml forbidden terms scan	HIGH

All findings are traceable to automated checks. Evidence sources reference specific crawl data, API specs, or code analysis results.