


# Coder

## Documentation Infrastructure Audit Report

Prepared by VeriOps

2026-05-05 11:09 UTC

CONFIDENTIAL | FOR INTERNAL USE ONLY



Risk: Moderate

---

### Key Metrics

Pages scanned: 820

Report type: Premium Executive Audit


Methodology: 5-pillar automated analysis + expert synthesis

### Headline Findings

- Broken internal links: 94
- Link verification inconclusive for 1 URLs (auth/rate-limit/server block); these are excluded from broken-link counts.
- SEO/GEO issue rate is high: 16.83%
- Freshness visibility is weak: only 38.29% pages show last updated

This report is generated automatically by the VeriOps platform. Estimates are directional.

## Executive Summary



Documentation audit of `coder.com/docs` covered 820 pages with 100% crawl coverage and 91.08% overall confidence. Critical issue: 94 confirmed broken internal links exist within user-facing documentation. API reference coverage is 75% with 276 API pages detected. Freshness metadata is visible on only 38.29% of pages, limiting maintainability transparency. SEO/GEO issues affect 16.83% of content.

External posture: SEO/GEO issue rate **16.8%**, public API coverage **75.0%**, broken links **94**.

|                             |                           |                                |
|-----------------------------|---------------------------|--------------------------------|
| <b>83</b><br>Audit Score    | <b>B+</b><br>Grade        | <b>Moderate</b><br>Risk Band   |
| <b>820</b><br>Pages Scanned | <b>94</b><br>Broken Links | <b>16.8%</b><br>SEO Issue Rate |

### Per-Site Breakdown


| Site URL                            | Pages | Broken Links | API Cov % | Example Rel % | SEO Issue % |
|-------------------------------------|-------|--------------|-----------|---------------|-------------|
| <code>https://coder.com/docs</code> | 820   | 94           | 75.0      | 88.6          | 16.8        |

## Industry Benchmark Comparison


| Benchmark                         | Score      | Gap        |
|-----------------------------------|------------|------------|
| Top performers (Stripe, Twilio)   | <b>92+</b> | <b>-9</b>  |
| Industry average (developer docs) | <b>85</b>  | <b>-2</b>  |
| Minimum acceptable                | <b>70</b>  | <b>+13</b> |
| <b>Your score</b>                 | <b>83</b>  | <b>-</b>   |

Gap to industry average: **2 points**. Top-performing developer documentation platforms (Stripe, Twilio, Vercel) score 92+. Closing this gap reduces support ticket volume by an estimated 20-30%.

## Board-Level Metrics



### Internal Metrics (from scorecard)



## Financial Exposure Model

| Low Estimate                                                                                  | Base Estimate                                                                                  | High Estimate                                                                                  |
|-----------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------|
| <b>\$140,822</b><br>annual exposure                                                           | <b>\$249,513</b><br>annual exposure                                                            | <b>\$431,728</b><br>annual exposure                                                            |
| Remediation: <b>\$2,090</b><br>Monthly loss: <b>\$4,125</b><br>Opportunity: <b>\$7,436/mo</b> | Remediation: <b>\$5,445</b><br>Monthly loss: <b>\$12,903</b><br>Opportunity: <b>\$7,436/mo</b> | Remediation: <b>\$8,800</b><br>Monthly loss: <b>\$27,808</b><br>Opportunity: <b>\$7,436/mo</b> |

Remediation: finding-level effort model | Monthly loss: operational friction model | Opportunity: engineering/support/release delay

Risk Matrix

| ID                     | Issue                                  | Severity | Monthly Loss | Fix Cost |
|------------------------|----------------------------------------|----------|--------------|----------|
| F-LAYERS               | Missing required doc layers            | HIGH     | \$1,782      | \$825    |
|                        |                                        |          |              |          |
| F-FRESHNESS            | Documentation freshness debt           | MEDIUM   | \$1,122      | \$550    |
|                        |                                        |          |              |          |
| F-RETRIEVAL            | RAG retrieval quality risk             | MEDIUM   | \$2,257      | \$990    |
|                        |                                        |          |              |          |
| F-API-COVERAGE         | Undocumented API operations            | LOW      | \$2,904      | \$1,100  |
|                        |                                        |          |              |          |
| F-EXAMPLES-RELIABILITY | Code examples fail or are non-runnable | LOW      | \$2,244      | \$935    |
|                        |                                        |          |              |          |
| F-DRIFT                | Code/docs contract drift               | LOW      | \$1,901      | \$660    |
|                        |                                        |          |              |          |
| F-TERMINOLOGY          | Terminology inconsistency              | LOW      | \$693        | \$385    |
|                        |                                        |          |              |          |

Priority Action Plan (Next 14 Days)

| ID                     | Finding                                | Severity |
|------------------------|----------------------------------------|----------|
| F-LAYERS               | Missing required doc layers            | HIGH     |
|                        |                                        |          |
| F-FRESHNESS            | Documentation freshness debt           | MEDIUM   |
|                        |                                        |          |
| F-RETRIEVAL            | RAG retrieval quality risk             | MEDIUM   |
|                        |                                        |          |
| F-API-COVERAGE         | Undocumented API operations            | LOW      |
|                        |                                        |          |
| F-EXAMPLES-RELIABILITY | Code examples fail or are non-runnable | LOW      |
|                        |                                        |          |

Recommended Actions

1. Repair all 94 confirmed broken internal links in user-facing documentation [impact: critical, effort: medium]
2. Implement last-updated metadata on remaining 61.71% of pages lacking freshness indicators [impact: high, effort: low]

3. Investigate and close 25% API reference coverage gap (identify missing endpoints or improve detection) [impact: high, effort: medium]
4. Audit and resolve 16.83% SEO/GEO issues (meta descriptions, titles, structured data) [impact: medium, effort: medium]
5. Investigate 444-page gap between requested and discovered pages to identify redirect chains or scope leaks [impact: medium, effort: high]

## Expert Analysis: Strengths and Risks

### Strengths

- + Complete crawl coverage achieved: 820 of 821 discovered pages indexed (100%)
- + High overall audit confidence at 91.08% with strong link verification confidence at 93.95%
- + API documentation present and detected across 276 pages with 75% reference coverage
- + Example code reliability estimated at 88.62%, indicating generally trustworthy samples
- + No repository navigation link breakage (0 repo\_broken\_links\_count)

### Risks

- 94 confirmed broken internal links in user-facing documentation create poor user experience and erode trust
- 38.29% freshness coverage means 61.71% of pages lack visible last-updated dates, obscuring content staleness
- 16.83% SEO/GEO issue rate may impact discoverability and search ranking
- 25% API reference gap leaves portion of API surface undocumented or undetected
- 1 unverified link due to auth/rate-limit/server blocks introduces uncertainty in link health
- 444 pages requested but not crawled (1265 requested vs 821 discovered) suggests potential scope or redirect issues
- Multi-project topology mode may hide inconsistencies across product documentation sets

### Limitations

- \* External audit cannot verify accuracy or correctness of technical content, code examples, or API specifications
- \* 1 unverified link due to authentication or rate-limiting means link health for protected resources is unconfirmed
- \* API coverage detection may miss non-standard documentation patterns or dynamically generated reference pages
- \* Freshness metadata reflects only visible last-updated dates; actual content currency requires manual review
- \* Multi-project topology may obscure cross-project linking issues or inconsistent standards between product doc sets

## Methodology

This audit applies a rigorous, repeatable methodology combining automated analysis with expert synthesis. The platform evaluates documentation quality across five pillars.

### 1. Automated Crawl + Structural HTML Analysis

Every public documentation page is crawled and analyzed for structural integrity, link health, navigation depth, and content organization. The crawler detects broken links, orphaned pages, redirect chains, and missing metadata.

### 2. SEO/GEO Scoring Model (24 Checks)

Each page is evaluated against 8 GEO checks (LLM/AI search optimization) and 16 SEO checks (traditional search engine optimization). Checks cover meta descriptions, heading hierarchy, fact density, URL structure, internal linking, and content depth.

### 3. API Reference Cross-Coverage Analysis

The platform compares published API reference documentation against the actual API surface (OpenAPI specs, GraphQL schemas, gRPC protos). Coverage gaps, undocumented endpoints, and stale references are identified.

### 4. Code Example Reliability Estimation

Code examples embedded in documentation are analyzed for syntactic correctness, completeness, and consistency with current API signatures. Reliability scores estimate the likelihood of examples working as documented.

### 5. Executive Synthesis and Prioritization

All quantitative findings are synthesized into actionable executive narratives. The analysis identifies patterns, prioritizes risks, and generates strategic recommendations calibrated to business impact.

Appendix A -- Economic Model Assumptions

| Assumption                          | Value |
|-------------------------------------|-------|
| engineer_hourly_usd                 | 110.0 |
| support_hourly_usd                  | 50.0  |
| release_count_per_month             | 8.0   |
| baseline_manual_sync_hours_per_week | 12.0  |
| avg_release_delay_hours             | 4.5   |
| monthly_support_tickets             | 800.0 |
| docs_related_ticket_share           | 0.22  |
| avg_ticket_handling_minutes         | 24.0  |

Note: estimates are directional and should be calibrated with client rates, release cadence, and support volume.

Appendix B -- Evidence Traceability

| Finding                                                          | Evidence Source                                   | Confidence |
|------------------------------------------------------------------|---------------------------------------------------|------------|
| F-LAYERS -- Missing required doc layers                          | policy pack + frontmatter layer scan              | HIGH       |
| F-FRESHNESS -- Documentation freshness debt                      | frontmatter date scan                             | HIGH       |
| F-RETRIEVAL -- RAG retrieval quality risk                        | retrieval_evals_report.json                       | HIGH       |
| F-API-COVERAGE -- Undocumented API operations                    | OpenAPI + docs text scan                          | HIGH       |
| F-EXAMPLES-RELIABILITY -- Code examples fail or are non-runnable | examples_smoke_report.json                        | HIGH       |
| F-DRIFT -- Code/docs contract drift                              | pr_docs_contract.json + api_sdk_drift_report.json | HIGH       |
| F-TERMINOLOGY -- Terminology inconsistency                       | glossary.yml forbidden terms scan                 | HIGH       |

All findings are traceable to automated checks. Evidence sources reference specific crawl data, API specs, or code analysis results.