


Turso

Documentation Infrastructure Audit Report

Prepared by VeriOps

2026-06-17 10:05 UTC

CONFIDENTIAL | FOR INTERNAL USE ONLY



Risk: High

Key Metrics

Pages scanned: 266

Report type: Premium Executive Audit


Methodology: 5-pillar automated analysis + expert synthesis

Headline Findings

- Broken internal links: 4
- SEO/GEO issue rate is high: 17.11%
- API usage-doc coverage is low: 43.75% (168 API pages detected)
- Example reliability estimate is low: 0.0%

This report is generated automatically by the VeriOps platform. Estimates are directional.

Executive Summary



Turso documentation audit covered 266 pages with 100.38% crawl coverage across a multi-project topology. Four confirmed broken internal links exist, all within user-facing docs. API reference coverage is 43.75% across 168 detected API pages, indicating significant undocumented surface area. Example reliability is 0.0%, freshness metadata appears on only 5.64% of pages, and 17.11% of pages have SEO/geo issues. Overall retrieval readiness scores 76.65 (moderate band) with 73.93% stale risk due to missing freshness signals.

External posture: SEO/GEO issue rate **17.1%**, public API coverage **43.8%**, broken links **4**.

64 Audit Score	C Grade	High Risk Band
266 Pages Scanned	4 Broken Links	17.1% SEO Issue Rate

Per-Site Breakdown


Site URL	Pages	Broken Links	API Cov %	Example Rel %	SEO Issue %
https://docs.turso.tech/	266	4	43.8	0.0	17.1

Industry Benchmark Comparison


Benchmark	Score	Gap
Top performers (Stripe, Twilio)	92+	-28
Industry average (developer docs)	85	-22
Minimum acceptable	70	-6
Your score	64	-

Gap to industry average: **22 points**. Top-performing developer documentation platforms (Stripe, Twilio, Vercel) score 92+. Closing this gap reduces support ticket volume by an estimated 20-30%.

Board-Level Metrics





Internal Metrics (from scorecard)




Pipeline Opportunity Fit (Deterministic)

This section maps cold-audit signals to problem classes that the pipeline can remediate directly. No LLM inference is used in this mapping.





LLM retrieval readiness breakdown (deterministic)



Evidence coverage: 97.3% (644/662 key claims backed by deterministic evidence signals).

Highest-value remediation targets

Problem	Severity	Pipeline solution path
Internal documentation link integrity issues	MEDIUM	Weekly link governance via automated audits and consolidated remediation queue
Search and AI readability quality gaps	HIGH	Deterministic SEO/GEO linting and fix cycles before publish
Code example reliability risk	HIGH	Snippet lint + smoke checks + protocol self-verify coverage
Weak freshness and lifecycle visibility	MEDIUM	Lifecycle management, stale detection, and weekly governance cadence
API reference coverage and drift risk	HIGH	API-first contract flow, validators, regression gates, and drift checks

Financial Exposure Model

Low Estimate \$165,619 annual exposure Remediation: \$1,615 Monthly loss: \$3,188 Opportunity: \$10,480/mo	Base Estimate \$249,608 annual exposure Remediation: \$4,208 Monthly loss: \$9,970 Opportunity: \$10,480/mo	High Estimate \$390,410 annual exposure Remediation: \$6,800 Monthly loss: \$21,488 Opportunity: \$10,480/mo
--	---	--

Remediation: finding-level effort model | Monthly loss: operational friction model | Opportunity: engineering/support/release delay

Risk Matrix

ID	Issue	Severity	Monthly Loss	Fix Cost
F-API-COVERAGE	Undocumented API operations	HIGH	\$2,244	\$850
F-EXAMPLES-RELIABILITY	Code examples fail or are non-runnable	HIGH	\$1,734	\$722
F-FRESHNESS	Documentation freshness debt	HIGH	\$867	\$425
F-LAYERS	Missing required doc layers	HIGH	\$1,377	\$638
F-DRIFT	Code/docs contract drift	LOW	\$1,469	\$510
F-TERMINOLOGY	Terminology inconsistency	LOW	\$536	\$298
F-RETRIEVAL	RAG retrieval quality risk	LOW	\$1,744	\$765

Priority Action Plan (Next 14 Days)

ID	Finding	Severity
F-API-COVERAGE	Undocumented API operations	HIGH
F-EXAMPLES-RELIABILITY	Code examples fail or are non-runnable	HIGH
F-FRESHNESS	Documentation freshness debt	HIGH
F-LAYERS	Missing required doc layers	HIGH
F-DRIFT	Code/docs contract drift	LOW

Recommended Actions

1. Add last-updated timestamps to all 266 pages via CI/CD automation or frontmatter injection [impact: high, effort: low]
2. Document the 94+ missing API endpoints (56.25% gap) with usage examples and parameter tables [impact: critical, effort: high]

3. Fix all 4 confirmed broken internal links in user-facing docs [impact: medium, effort: low]
4. Implement tested, runnable examples for top 20 API pages to lift 0.0% example reliability baseline [impact: high, effort: medium]
5. Resolve 45+ SEO/geo issues (17.11% rate) via meta tag audits, canonical URL fixes, and hreflang setup [impact: medium, effort: medium]

Expert Analysis: Strengths and Risks

Strengths

- + Complete crawl coverage (100.38%) with 266 pages discovered and indexed
- + High evidence coverage at 97.28% (644 of 662 claims backed by evidence)
- + Perfect answerability and citationability scores (100.0%) for AI retrieval
- + No repository navigation broken links; all 4 broken links are in user-facing docs
- + Strong overall confidence at 91.5% with reliable detection across metrics

Risks

- 56.25% of API surface (94+ API pages) lacks usage documentation, blocking developer onboarding and integration success
- 0.0% example reliability means no runnable/tested examples detected, risking user frustration and support load
- 73.93% stale risk due to 94.36% of pages missing last-updated metadata, eroding trust and AI answer quality
- 17.11% SEO/geo issue rate (45+ pages) limits discoverability in search and regional access
- 4 confirmed broken internal links create dead ends in user navigation and damage credibility
- Only 66.67% actionability coverage suggests many pages lack clear next steps for users
- 43.75% API reference coverage may reflect incomplete API catalog or missing auto-generation pipeline

Limitations

- * External audit cannot verify accuracy of technical claims, code correctness, or whether examples execute successfully in live environments.
- * API coverage detection relies on heuristic page classification; actual API surface may differ if endpoints are documented in non-standard formats or external systems.
- * Freshness and stale risk metrics depend on metadata presence; pages may be current but lack timestamps, or timestamps may not reflect content accuracy.
- * SEO/geo issue detection covers technical tags and structure but cannot assess content quality, keyword strategy, or competitive search positioning.
- * Actionability and example reliability scores reflect structural patterns, not user task completion rates or real-world developer success metrics.

Methodology

This audit applies a rigorous, repeatable methodology combining automated analysis with expert synthesis. The platform evaluates documentation quality across five pillars.

1. Automated Crawl + Structural HTML Analysis

Every public documentation page is crawled and analyzed for structural integrity, link health, navigation depth, and content organization. The crawler detects broken links, orphaned pages, redirect chains, and missing metadata.

2. SEO/GEO Scoring Model (24 Checks)

Each page is evaluated against 8 GEO checks (LLM/AI search optimization) and 16 SEO checks (traditional search engine optimization). Checks cover meta descriptions, heading hierarchy, fact density, URL structure, internal linking, and content depth.

3. API Reference Cross-Coverage Analysis

The platform compares published API reference documentation against the actual API surface (OpenAPI specs, GraphQL schemas, gRPC protos). Coverage gaps, undocumented endpoints, and stale references are identified.

4. Code Example Reliability Estimation

Code examples embedded in documentation are analyzed for syntactic correctness, completeness, and consistency with current API signatures. Reliability scores estimate the likelihood of examples working as documented.

5. Executive Synthesis and Prioritization

All quantitative findings are synthesized into actionable executive narratives. The analysis identifies patterns, prioritizes risks, and generates strategic recommendations calibrated to business impact.

Appendix A -- Economic Model Assumptions

Assumption	Value
engineer_hourly_usd	85.0
support_hourly_usd	55.0
release_count_per_month	8.0
baseline_manual_sync_hours_per_week	4.0
avg_release_delay_hours	18.0
monthly_support_tickets	180.0
docs_related_ticket_share	0.28
avg_ticket_handling_minutes	22.0
monthly_doc_influenced_evals	45.0
eval_to_customer_rate	0.12
avg_customer_monthly_value_usd	850.0
docs_friction_revenue_sensitivity	0.35

Note: estimates are directional and should be calibrated with client rates, release cadence, and support volume.

Appendix B -- Evidence Traceability

Finding	Evidence Source	Confidence
F-API-COVERAGE -- Undocumented API operations	OpenAPI + docs text scan	HIGH
F-EXAMPLES-RELIABILITY -- Code examples fail or are non-runnable	examples_smoke_report.json	HIGH
F-FRESHNESS -- Documentation freshness debt	frontmatter date scan	HIGH
F-LAYERS -- Missing required doc layers	policy pack + frontmatter layer scan	HIGH
F-DRIFT -- Code/docs contract drift	pr_docs_contract.json + api_sdk_drift_report.json	HIGH
F-TERMINOLOGY -- Terminology inconsistency	glossary.yml forbidden terms scan	LOW

F-RETRIEVAL -- RAG retrieval quality risk	retrieval_evals_report.json	HIGH
---	-----------------------------	------

All findings are traceable to automated checks. Evidence sources reference specific crawl data, API specs, or code analysis results.